

House Prices - Advanced Regression Techniques

Background & Objectives

Ask a home buyer to describe their dream house, and they probably won't begin with the height of the basement ceiling or the proximity to an east-west railroad. The dataset proves that much more influences price negotiations than the number of bedrooms or a white-picket fence.

Objective: With 79 explanatory variables describing (almost) every aspect of residential homes in Ames, Iowa, the project predict the final price of each home.

In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from scipy.stats import norm
from sklearn.preprocessing import StandardScaler
from scipy import stats
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import ElasticNet, Lasso, BayesianRidge, LassoLarsIC
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.kernel_ridge import KernelRidge
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import RobustScaler
from sklearn.base import BaseEstimator, TransformerMixin, RegressorMixin, clone
from sklearn.model_selection import KFold, cross_val_score, train_test_split
from sklearn.metrics import mean_squared_error
import xgboost as xgb
import lightgbm as lgb
```

Data Import & Description

In [2]:

```
train_csv = pd.read_csv('train.csv')
test_csv = pd.read_csv('test.csv')
```

In [3]:

```
train_csv['dataset'] = '1'
test_csv['dataset'] = '0'
hp_csv = pd.concat([train_csv, test_csv], axis=0)
```

In []:

```
train_csv.
```

In [4]:

```
train_csv.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1460 entries, 0 to 1459
```

```
Data columns (total 82 columns):
```

#	Column	Non-Null Count	Dtype
0	Id	1460 non-null	int64
1	MSSubClass	1460 non-null	int64
2	MSZoning	1460 non-null	object
3	LotFrontage	1201 non-null	float64
4	LotArea	1460 non-null	int64
5	Street	1460 non-null	object
6	Alley	91 non-null	object
7	LotShape	1460 non-null	object
8	LandContour	1460 non-null	object
9	Utilities	1460 non-null	object
10	LotConfig	1460 non-null	object
11	LandSlope	1460 non-null	object
12	Neighborhood	1460 non-null	object
13	Condition1	1460 non-null	object
14	Condition2	1460 non-null	object

In [5]:

test_csv.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1459 entries, 0 to 1458

Data columns (total 81 columns):

#	Column	Non-Null Count	Dtype
0	Id	1459 non-null	int64
1	MSSubClass	1459 non-null	int64
2	MSZoning	1455 non-null	object
3	LotFrontage	1232 non-null	float64
4	LotArea	1459 non-null	int64
5	Street	1459 non-null	object
6	Alley	107 non-null	object
7	LotShape	1459 non-null	object
8	LandContour	1459 non-null	object
9	Utilities	1457 non-null	object
10	LotConfig	1459 non-null	object
11	LandSlope	1459 non-null	object
12	Neighborhood	1459 non-null	object
13	Condition1	1459 non-null	object
14	Condition2	1459 non-null	object
15	BldgType	1459 non-null	object
16	HouseStyle	1459 non-null	object
17	OverallQual	1459 non-null	int64
18	OverallCond	1459 non-null	int64
19	YearBuilt	1459 non-null	int64
20	YearRemodAdd	1459 non-null	int64
21	RoofStyle	1459 non-null	object
22	RoofMatl	1459 non-null	object
23	Exterior1st	1458 non-null	object
24	Exterior2nd	1458 non-null	object
25	MasVnrType	1443 non-null	object
26	MasVnrArea	1444 non-null	float64
27	ExterQual	1459 non-null	object
28	ExterCond	1459 non-null	object
29	Foundation	1459 non-null	object
30	BsmtQual	1415 non-null	object
31	BsmtCond	1414 non-null	object
32	BsmtExposure	1415 non-null	object
33	BsmtFinType1	1417 non-null	object
34	BsmtFinSF1	1458 non-null	float64
35	BsmtFinType2	1417 non-null	object
36	BsmtFinSF2	1458 non-null	float64
37	BsmtUnfSF	1458 non-null	float64
38	TotalBsmtSF	1458 non-null	float64
39	Heating	1459 non-null	object
40	HeatingQC	1459 non-null	object
41	CentralAir	1459 non-null	object
42	Electrical	1459 non-null	object
43	1stFlrSF	1459 non-null	int64
44	2ndFlrSF	1459 non-null	int64
45	LowQualFinSF	1459 non-null	int64
46	GrLivArea	1459 non-null	int64
47	BsmtFullBath	1457 non-null	float64
48	BsmtHalfBath	1457 non-null	float64
49	FullBath	1459 non-null	int64
50	HalfBath	1459 non-null	int64
51	BedroomAbvGr	1459 non-null	int64

```

52 KitchenAbvGr      1459 non-null    int64
53 KitchenQual       1458 non-null    object
54 TotRmsAbvGrd      1459 non-null    int64
55 Functional         1457 non-null    object
56 Fireplaces        1459 non-null    int64
57 FireplaceQu       729 non-null     object
58 GarageType        1383 non-null    object
59 GarageYrBlt       1381 non-null    float64
60 GarageFinish      1381 non-null    object
61 GarageCars        1458 non-null    float64
62 GarageArea        1458 non-null    float64
63 GarageQual        1381 non-null    object
64 GarageCond        1381 non-null    object
65 PavedDrive        1459 non-null    object
66 WoodDeckSF        1459 non-null    int64
67 OpenPorchSF       1459 non-null    int64
68 EnclosedPorch     1459 non-null    int64
69 3SsnPorch         1459 non-null    int64
70 ScreenPorch       1459 non-null    int64
71 PoolArea          1459 non-null    int64
72 PoolQC            3 non-null       object
73 Fence             290 non-null     object
74 MiscFeature       51 non-null      object
75 MiscVal           1459 non-null    int64
76 MoSold            1459 non-null    int64
77 YrSold            1459 non-null    int64
78 SaleType          1458 non-null    object
79 SaleCondition     1459 non-null    object
80 dataset           1459 non-null    object

```

dtypes: float64(11), int64(26), object(44)

memory usage: 923.4+ KB

In [6]:

```
hp_csv.head()
```

Out[6]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Ut
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	A
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	A
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	A
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	A
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	A

5 rows × 82 columns

In [7]:

```
y=hp_csv['SalePrice'][hp_csv['dataset']=='1']  
y.describe()
```

Out[7]:


```
count      1460.000000  
mean       180921.195890  
std        79442.502883  
min         34900.000000  
25%        129975.000000  
50%        163000.000000  
75%        214000.000000  
max        755000.000000  
Name: SalePrice, dtype: float64
```

In [8]:

```
sns.distplot(y)
```

Out[8]:

<AxesSubplot:xlabel='SalePrice', ylabel='Density'>



Data processing

missing data

In [9]:

```
total=hp_csv.isnull().sum().sort_values(ascending=False)  
percent=(hp_csv.isnull().sum()/hp_csv.isnull().count()).sort_values(ascending=True)  
missingdata=pd.concat([total,percent],axis=1,keys=['total','percent'])  
missingdata=missingdata[missingdata['total']>0]
```

In [10]:

```
print(missingdata)
```

	total	percent
PoolQC	2909	0.996574
MiscFeature	2814	0.964029
Alley	2721	0.932169
Fence	2348	0.804385
SalePrice	1459	0.499829
FireplaceQu	1420	0.486468
LotFrontage	486	0.166495
GarageYrBlt	159	0.054471
GarageQual	159	0.054471
GarageFinish	159	0.054471
GarageCond	159	0.054471
GarageType	157	0.053786
BsmtCond	82	0.028092
BsmtExposure	82	0.028092
BsmtQual	81	0.027749
BsmtFinType2	80	0.027407
BsmtFinType1	79	0.027064
MasVnrType	24	0.008222

In [11]:

```
missd_name=list(missingdata.index)
object_list=[]
non_object_list=[]
for i in missd_name:
    if hp_csv[i].dtype==object:
        object_list.append(i)
    else:
        non_object_list.append(i)
print("object_list:",object_list)
print("non_object_list:",non_object_list)
```

```
object_list: ['PoolQC', 'MiscFeature', 'Alley', 'Fence', 'FireplaceQu',
'GarageQual', 'GarageFinish', 'GarageCond', 'GarageType', 'BsmtCond',
'BsmtExposure', 'BsmtQual', 'BsmtFinType2', 'BsmtFinType1', 'MasVnrType',
'MSZoning', 'Functional', 'Utilities', 'Electrical', 'Exterior2nd',
'Exterior1st', 'SaleType', 'KitchenQual']
non_object_list: ['SalePrice', 'LotFrontage', 'GarageYrBlt', 'MasVnrArea',
'BsmtFullBath', 'BsmtHalfBath', 'BsmtFinSF1', 'GarageCars', 'GarageArea',
'TotalBsmtSF', 'BsmtUnfSF', 'BsmtFinSF2']
```

In [12]:

```
for i in object_list:
    hp_csv[i].fillna('None',inplace=True)
```

In [13]:

```
for i in non_object_list[2:]:
    hp_csv[i].fillna(0,inplace=True)
```

For those over 40% missing values, it means that it can be deleted. Therefore, 'PoolQC', 'MiscFeature', 'Alley', 'Fence', and 'FireplaceQu' should be deleted.

In [14]:

```
hp_new=hp_csv.drop(['PoolQC', 'MiscFeature', 'Alley', 'Fence', 'FireplaceQu'],axis=1)
```

In [15]:

```
hp_csv['LotFrontage']=hp_csv.groupby(['Neighborhood']).transform(lambda x: x.fillna(
```

In [16]:

```
hp_csv['LotFrontage']=hp_csv['LotFrontage'].astype(int)
```

Variable processing and encoding

In [17]:

```
full_name=list(hp_csv.columns)
object_full_list=[]
nonobject_full_list=[]
for i in full_name:
    if hp_csv[i].dtype==object:
        object_full_list.append(i)
    else:
        nonobject_full_list.append(i)
print("object_list:",object_full_list)
print("nonobject_list:",nonobject_full_list)
```

```
object_list: ['MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2', 'Heating', 'HeatingQC', 'CentralAir', 'Electrical', 'KitchenQual', 'Functional', 'FireplaceQu', 'GarageType', 'GarageFinish', 'GarageQual', 'GarageCond', 'PavedDrive', 'PoolQC', 'Fence', 'MiscFeature', 'SaleType', 'SaleCondition', 'dataset']
nonobject_list: ['Id', 'MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces', 'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal', 'MoSold', 'YrSold', 'SalePrice']
```

In [18]:

```
for i in object_full_list+non_object_list:
    lbl=LabelEncoder()
    lbl.fit(list(hp_csv[i].values))
    hp_csv[i]=lbl.transform(list(hp_csv[i].values))
```

In [19]:

```
hp_csv.head()
```

Out[19]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Ut
0	1	60	4	0	8450	1	1	3	3	
1	2	20	4	1	9600	1	1	3	3	
2	3	60	4	2	11250	1	1	0	3	
3	4	70	4	3	9550	1	1	0	3	
4	5	60	4	4	14260	1	1	0	3	

5 rows × 82 columns

In [20]:

```
from sklearn.preprocessing import MinMaxScaler
hp_new2=hp_csv.drop(['Id', 'SalePrice'],axis=1)
col_name=list(hp_new2.columns)
scaler=MinMaxScaler().fit(hp_new2[col_name])
hp_new2[col_name]=scaler.transform(hp_new2[col_name])
hp_new2
```

Out[20]:

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Ut
0	0.235294	0.8	0.000000	0.033420	1.0	0.5	1.0	1.0	
1	0.000000	0.8	0.000343	0.038795	1.0	0.5	1.0	1.0	
2	0.235294	0.8	0.000685	0.046507	1.0	0.5	0.0	1.0	
3	0.294118	0.8	0.001028	0.038561	1.0	0.5	0.0	1.0	
4	0.235294	0.8	0.001371	0.060576	1.0	0.5	0.0	1.0	
...
1454	0.823529	1.0	0.998629	0.002973	1.0	0.5	1.0	1.0	
1455	0.823529	1.0	0.998972	0.002776	1.0	0.5	1.0	1.0	
1456	0.000000	0.8	0.999315	0.087406	1.0	0.5	1.0	1.0	
1457	0.382353	0.8	0.999657	0.042726	1.0	0.5	1.0	1.0	
1458	0.235294	0.8	1.000000	0.038921	1.0	0.5	1.0	1.0	

2919 rows × 80 columns

In [21]:

```

full_col=list(hp_new2.columns)
train_cc=hp_new2[full_col]
skew_cc=train_cc.skew(axis=0)
skew=skew_cc[abs(skew_cc)>0.75]
skew=skew.to_frame(name='skewness')
print(skew)
print(skew.index)

```

	skewness
MSSubClass	1.376165
MSZoning	-2.575195
LotArea	12.829025
Street	-15.508104
LandContour	-3.118295
Utilities	33.990952
LotConfig	-1.196901
LandSlope	4.977715
Condition1	2.984648
Condition2	12.066294
BldgType	2.193388
RoofStyle	1.554106
RoofMatl	8.712245
MasVnrArea	1.557110
ExterQual	-1.802335
ExterCond	-2.499003
BsmtCond	-2.864057
BsmtExposure	-1.115430
BsmtFinType2	-3.045893
BsmtFinSF2	3.478350
Heating	12.084999
CentralAir	-3.460801
Electrical	-3.049382
1stFlrSF	1.470360
2ndFlrSF	0.862118
LowQualFinSF	12.094977
GrLivArea	1.270010
BsmtHalfBath	3.933616
KitchenAbvGr	4.304467
TotRmsAbvGrd	0.758757
Functional	-3.820689
GarageQual	-3.075732
GarageCond	-3.597639
PavedDrive	-2.980616
WoodDeckSF	1.843380
OpenPorchSF	2.536417
EnclosedPorch	4.005950
3SsnPorch	11.381914
ScreenPorch	3.948723
PoolArea	16.907017
PoolQC	-20.734650
Fence	-1.994802
MiscFeature	5.066925
MiscVal	21.958480
SaleType	-3.325081
SaleCondition	-2.789472
Index(['MSSubClass', 'MSZoning', 'LotArea', 'Street', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope', 'Condition1', 'Condition 2', 'BldgType', 'RoofStyle', 'RoofMatl', 'MasVnrArea', 'ExterQual',	

```

    'ExterCond', 'BsmtCond', 'BsmtExposure', 'BsmtFinType2', 'BsmtFinSF2',
    'Heating', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
    'LowQualFinSF', 'GrLivArea', 'BsmtHalfBath', 'KitchenAbvGr',
    'TotRmsAbvGrd', 'Functional', 'GarageQual', 'GarageCond', 'PavedDrive',
    'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
    'ScreenPorch', 'PoolArea', 'PoolQC', 'Fence', 'MiscFeature', 'MiscVal',
    'SaleType', 'SaleCondition'],
    dtype='object')

```

In [22]:

```

for i in list(skew.index):
    hp_csv[i]=np.log1p(hp_csv[i])

```

In [23]:

```

full_col=list(hp_csv.columns)
train_cc=hp_csv[full_col]
skew_cc=train_cc.skew(axis=0)
skew=skew_cc[abs(skew_cc)>0.75]
skew=skew.to_frame(name='skewness')
print(skew)
print(skew.index)

```

	skewness
MSZoning	-3.749173
Street	-15.508104
LandContour	-3.528663
Utilities	32.722387
LotConfig	-1.353095
LandSlope	4.486542
Condition1	-0.959142
Condition2	-4.781968
BldgType	1.971066
RoofStyle	1.294681
RoofMatl	8.254887
ExterQual	-3.049371
ExterCond	-3.182679
BsmtCond	-3.320867
BsmtExposure	-1.457024
BsmtFinType2	-3.951148
BsmtFinSF2	2.515870
Heating	9.592262

In [24]:

```
hp_new2.describe()
```

Out[24]:

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotSh
count	2919.000000	2919.000000	2919.000000	2919.000000	2919.000000	2919.000000	2919.000
mean	0.218457	0.793765	0.500000	0.041450	0.995889	0.492806	0.649
std	0.250104	0.171681	0.288824	0.036865	0.063996	0.130046	0.469
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000
25%	0.000000	0.800000	0.250000	0.028877	1.000000	0.500000	0.000
50%	0.176471	0.800000	0.500000	0.038108	1.000000	0.500000	1.000
75%	0.294118	0.800000	0.750000	0.048003	1.000000	0.500000	1.000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000

8 rows × 80 columns

In [89]:

```
train_new=hp_new2[hp_new2['dataset']==1]
test_new=hp_new2[hp_new2['dataset']==0]
```

feature selection and deduction

K-means

In [27]:

```
train_real=train_new.drop(['dataset'],axis=1)
```

In [90]:


```
test_real = test_new.drop(['dataset'],axis=1)
```

In [30]:

```
from sklearn.cluster import KMeans
inertia=[]
list_num_clusters = list(range(1,11))
for num_clusters in list_num_clusters:
    km=KMeans(n_clusters=num_clusters)
    km.fit(train_real)
    inertia.append(km.inertia_)
```

In [31]:

```
plt.plot(list_num_clusters,inertia)
plt.scatter(list_num_clusters,inertia)
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.show()
```



In [32]:

```
km2=KMeans(n_clusters=3).fit(train_real)
train_real['group']=km2.labels_
```

In [33]:

```
train_new2=pd.concat([train_new,train_real['group']],axis=1)
```

In [34]:

```
train_new2.head()
```

Out[34]:

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities
0	0.235294	0.8	0.000000	0.033420	1.0	0.5	1.0	1.0	0
1	0.000000	0.8	0.000343	0.038795	1.0	0.5	1.0	1.0	0
2	0.235294	0.8	0.000685	0.046507	1.0	0.5	0.0	1.0	0
3	0.294118	0.8	0.001028	0.038561	1.0	0.5	0.0	1.0	0
4	0.235294	0.8	0.001371	0.060576	1.0	0.5	0.0	1.0	0

5 rows × 81 columns

In [43]:

```
from pandas import DataFrame  
y_sale=DataFrame({'SalePrice':y})
```

In [44]:

```
y_sale
```

Out[44]:

	SalePrice
0	208500.0
1	181500.0
2	223500.0
3	140000.0
4	250000.0
...	...
1455	175000.0
1456	210000.0
1457	266500.0
1458	142125.0
1459	147500.0

1460 rows × 1 columns

In [46]:

```
train_new3=pd.concat([train_new2,y_sale],axis=1)
print(train_new3.head())
```

```

   MSSubClass  MSZoning  LotFrontage  LotArea  Street  Alley  LotShap
e \
0    0.235294      0.8      0.000000  0.033420      1.0    0.5      1.
0
1    0.000000      0.8      0.000343  0.038795      1.0    0.5      1.
0
2    0.235294      0.8      0.000685  0.046507      1.0    0.5      0.
0
3    0.294118      0.8      0.001028  0.038561      1.0    0.5      0.
0
4    0.235294      0.8      0.001371  0.060576      1.0    0.5      0.
0

   LandContour  Utilities  LotConfig  ...  Fence  MiscFeature  MiscVal
\
0           1.0         0.0         1.0  ...    1.0           0.25      0.0
1           1.0         0.0         0.5  ...    1.0           0.25      0.0
2           1.0         0.0         1.0  ...    1.0           0.25      0.0
3           1.0         0.0         0.0  ...    1.0           0.25      0.0
4           1.0         0.0         0.5  ...    1.0           0.25      0.0

   MoSold  YrSold  SaleType  SaleCondition  dataset  group  SalePric
e
0  0.090909   0.50        1.0             0.8      1.0      1  208500.
0
1  0.363636   0.25        1.0             0.8      1.0      0  181500.
0
2  0.727273   0.50        1.0             0.8      1.0      1  223500.
0
3  0.090909   0.00        1.0             0.0      1.0      0  140000.
0
4  1.000000   0.50        1.0             0.8      1.0      1  250000.
0

[5 rows x 82 columns]
```

In [47]:

```
group_data=train_new3.groupby(['group']).mean()
group_data['SalePrice']
```

Out[47]:

```

group
0    152732.001890
1    240521.496610
2    121530.891496
Name: SalePrice, dtype: float64
```

Difference in sales is obvious so it may be a good clustering. However, too many dimensions end in 3 clusters might lose a lot of information. The target is to predict the housing price, so the K-means method is not applicable for predicting housing prices. Next Steps: (1) Used PCA to decrease the dimensions (2) Used models (Regression, Bagging and boosting methods) to predict the housing prices

PCA

In [48]:

```
train_real2=train_real.drop('group',axis=1)
```

In [49]:

```
train_real2.shape
```

Out[49]:

```
(1460, 79)
```

In [91]:

```
test_real.shape
```

Out[91]:

```
(1459, 79)
```

In [51]:

```

from sklearn.decomposition import PCA
pca_list=list()
feature_weight_list=list()
for n in range(1,40):
    PCAmod=PCA(n_components=n)
    PCAmod.fit(train_real2)
    pca_list.append(pd.Series({'n':n, 'model':PCAmod,
                              'var':PCAmod.explained_variance_ratio_.sum()}))
    abs_feature_values=np.abs(PCAmod.components_).sum(axis=0)
    feature_weight_list.append(pd.DataFrame({
        'n':n,
        'features':train_real2.columns,
        'values':abs_feature_values/abs_feature_values.sum()
    }))
pca_df=pd.concat(pca_list,axis=1).T.set_index('n')
pca_df

```


Out[51]:

	model	var
n		
1	PCA(n_components=1)	0.176047
2	PCA(n_components=2)	0.253102
3	PCA(n_components=3)	0.315185
4	PCA(n_components=4)	0.36202
5	PCA(n_components=5)	0.402422
6	PCA(n_components=6)	0.440148
7	PCA(n_components=7)	0.475287
8	PCA(n_components=8)	0.507335
9	PCA(n_components=9)	0.538172
10	PCA(n_components=10)	0.566593
11	PCA(n_components=11)	0.592579
12	PCA(n_components=12)	0.614398
13	PCA(n_components=13)	0.634967
14	PCA(n_components=14)	0.654127
15	PCA(n_components=15)	0.672641
16	PCA(n_components=16)	0.68879
17	PCA(n_components=17)	0.705023
18	PCA(n_components=18)	0.720251
19	PCA(n_components=19)	0.735925
20	PCA(n_components=20)	0.751022
21	PCA(n_components=21)	0.765266
22	PCA(n_components=22)	0.77862
23	PCA(n_components=23)	0.791751

	model	var
n		
24	PCA(n_components=24)	0.804461
25	PCA(n_components=25)	0.816113
26	PCA(n_components=26)	0.827413
27	PCA(n_components=27)	0.838586
28	PCA(n_components=28)	0.849058
29	PCA(n_components=29)	0.858548
30	PCA(n_components=30)	0.867633
31	PCA(n_components=31)	0.876647
32	PCA(n_components=32)	0.884759
33	PCA(n_components=33)	0.892515
34	PCA(n_components=34)	0.89962
35	PCA(n_components=35)	0.906323
36	PCA(n_components=36)	0.912802
37	PCA(n_components=37)	0.918935
38	PCA(n_components=38)	0.924703
39	PCA(n_components=39)	0.930199

In [52]:

```
sns.set_context('talk')
plt.figure(figsize=(15, 8))
ax = pca_df['var'].plot(kind='bar')
ax.set(xlabel='Number of dimensions',
       ylabel='Percent explained variance',
       title='Explained variance vs Dimensions');
```



79 dimensions could be deducted to 35 dimensions, with variances more than 90.6%

In [53]:

```
pca=PCA(n_components=35)
X_scaled=pca.fit_transform(train_real2)
print(X_scaled)
print(X_scaled.shape)
```

```
[ [ 0.62139404 -0.09039312 -0.56881117 ... -0.04505445  0.08500836
    -0.13048469]
  [ 0.14268384  0.57654452 -0.32490471 ...  0.06578555 -0.05847321
    -0.14479061]
  [ 0.87281173  0.12776005  0.21598469 ... -0.0675353  0.00393354
    -0.0256786 ]
  ...
  [ 0.0727726  -0.30419109 -0.31129305 ...  0.10830725  0.38508199
    0.02175034]
  [-0.61378183  0.22254815 -0.40579911 ... -0.00675809  0.10073956
    0.26285867]
  [-0.24102832  0.61016369 -0.51363074 ... -0.00328618  0.12384418
    0.13034319]]
(1460, 35)
```

In [92]:

```
pca=PCA(n_components=35)
Test_scaled=pca.fit_transform(test_real)
print(Test_scaled.shape)
```

```
(1459, 35)
```

Modelling

We choose several models and use 5-folds cross-validation to evaluate these models. Models include:

LinearRegression Ridge Lasso Random Forrest Gradient Boosting Tree Support Vector Regression Linear
Support Vector Regression ExtraTreesRegressor XgBoost Adaboost GBDT

In [55]:

```
from sklearn.model_selection import cross_val_score, GridSearchCV, KFold
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor, ExtraTreesRegressor
from sklearn.svm import SVR, LinearSVR
from sklearn.linear_model import ElasticNet, SGDRegressor, BayesianRidge
from sklearn.kernel_ridge import KernelRidge
from xgboost import XGBRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.ensemble import GradientBoostingRegressor
```

In [62]:

```
y1=np.log1p(y)
```

In [58]:

```
X_scaled=pd.DataFrame(X_scaled)
```

In [59]:

```
X_scaled.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1460 entries, 0 to 1459
```

```
Data columns (total 35 columns):
```

#	Column	Non-Null Count	Dtype
0	0	1460 non-null	float64
1	1	1460 non-null	float64
2	2	1460 non-null	float64
3	3	1460 non-null	float64
4	4	1460 non-null	float64
5	5	1460 non-null	float64
6	6	1460 non-null	float64
7	7	1460 non-null	float64
8	8	1460 non-null	float64
9	9	1460 non-null	float64
10	10	1460 non-null	float64
11	11	1460 non-null	float64
12	12	1460 non-null	float64
13	13	1460 non-null	float64
14	14	1460 non-null	float64

In [60]:

```
models = {}
models['LR'] = LinearRegression()
models['Ridge'] = Ridge()
models['Lasso'] = Lasso(alpha=0.01,max_iter=10000)
models['RF'] = RandomForestRegressor()
models['SVR'] = SVR()
models['LR'] = LinearSVR()
models['Ela'] = ElasticNet(alpha=0.001,max_iter=10000)
models['SGD'] = SGDRegressor(max_iter=1000,tol=1e-3)
models['Bay'] = BayesianRidge()
models['KR'] = KernelRidge(alpha=0.6, kernel='polynomial', degree=2, coef0=2.5)
models['Extra'] = ExtraTreesRegressor()
models['XGB'] = XGBRegressor()
models['Ada'] = AdaBoostRegressor()
models['GBDT'] = GradientBoostingRegressor()
```

In [63]:

```

kf = KFold(10)
for model in models:
    cv_result = cross_val_score(models[model], X_scaled, y1, cv=kf, scoring="neg_mean_squared_error")
    print('%s model scoring %s, std %s' % (model,
                                           round(np.sqrt(-cv_result.mean()),4),
                                           cv_result.std()))

```

```

LR model scoring 0.1535, std 0.006914052523403825
Ridge model scoring 0.153, std 0.007007363047927657
Lasso model scoring 0.1956, std 0.008428107648325018
RF model scoring 0.1791, std 0.00619778429454758
SVR model scoring 0.1592, std 0.004861613521901746
Ela model scoring 0.1531, std 0.006963691626275504
SGD model scoring 0.1641, std 0.007061279691485609
Bay model scoring 0.153, std 0.006999910729121719
KR model scoring 0.1508, std 0.006837252265078653
Extra model scoring 0.1731, std 0.006562451583580815
XGB model scoring 0.1674, std 0.006237492230090017
Ada model scoring 0.2084, std 0.0066340950257160036
GBDT model scoring 0.1699, std 0.006902300916498008

```

Narrow down to better performance models: Adaboost RF,Extra Lasso

Next we do some hyperparameters tuning. First define a gridsearch method.

Lasso

In [73]:

```

grid=GridSearchCV(estimator=Lasso(),
                  param_grid=
                      {'alpha':[0.0001,0.00015,0.0002,0.0005,0.01],
                       'max_iter':[10000],
                       'normalize':[True,False]},
                  cv=kf,
                  scoring='neg_mean_squared_error')
grid_result=grid.fit(X_scaled,y1)
print('Lasso best_estimator is %s,scoring %s'
      %(grid_result.best_params_,
        round(np.sqrt(-1*grid_result.best_score_),4)
      ))

```

```

Lasso best_estimator is {'alpha': 0.0002, 'max_iter': 10000, 'normalize': False},scoring 0.1529

```

Bagging Models

In [74]:

```

bagging_models={'RandomForest':RandomForestRegressor(),
                'ExtraTree':ExtraTreesRegressor()}
bagging_params={'n_estimators':[10,50,100,200,500,800]}

```

In [75]:

```

kf=KFold(10)
for model in bagging_models:
    grid=GridSearchCV(estimator=bagging_models[model],
                      param_grid=bagging_params,
                      cv=kf,
                      scoring='neg_mean_squared_error')
    grid_result=grid.fit(X_scaled,y1)
    print('%s model best_estimator is %s, scoring %s'
          %(model,
            grid_result.best_params_,
            round(np.sqrt(-1*grid_result.best_score_),4)))

```

RandomForest model best_estimator is {'n_estimators': 100}, scoring 0.1773
 ExtraTree model best_estimator is {'n_estimators': 500}, scoring 0.1731

Boosting models

In [78]:

```

boosting_models={'Adaboost':AdaBoostRegressor()}
boosting_params={'n_estimators':[10,50,100,200,500,800],
                  'learning_rate':[0.005,0.01,0.1,0.5,1]}
kf=KFold(10)
for model in boosting_models:
    grid=GridSearchCV(estimator=boosting_models[model],
                      param_grid=boosting_params,
                      cv=kf,
                      scoring='neg_mean_squared_error')
    grid_result=grid.fit(X_scaled,y1)
    print('%s model best_estimator is %s, scoring %s'
          %(model,
            grid_result.best_params_,
            round(np.sqrt(-1*grid_result.best_score_),4)))

```

Adaboost model best_estimator is {'learning_rate': 1, 'n_estimators': 100}, scoring 0.2034

In [93]:

```

ad = AdaBoostRegressor(n_estimators = 100, learning_rate = 1)
ad .fit(X_scaled,y1)
y_predit= ad.predict(Test_scaled)

```

In [95]:

```
y_predit
```

Out[95]:

```
array([11.69792134, 11.9451477 , 12.28589869, ..., 11.8111241 ,
       11.80571857, 12.40334841])
```

In [96]:

```
y_sale=np.expml(y_predit)
```

In [97]:

```
y_sale.shape
```

Out[97]:

```
(1459,)
```

In [98]:

```
test_csv.shape
```

Out[98]:

```
(1459, 81)
```

Submission

In [99]:

```
submission=pd.DataFrame({  
    'ID':test_csv['Id'],  
    'SalePrice':y_sale  
})  
submission.to_csv('submission.csv',index=False)
```

In [100]:

```
ls= Lasso(alpha=0.01,max_iter=10000)  
ls. fit (X_scaled,y1)  
y_lasso= ls.predict(Test_scaled)  
y_lasso=np.expml(y_lasso)  
submission=pd.DataFrame({  
    'ID':test_csv['Id'],  
    'SalePrice':y_lasso  
})  
submission.to_csv('submission2.csv',index=False)
```

In [103]:

```
submission=pd.DataFrame({  
    'ID':test_csv['Id'],  
    'SalePrice':(y_sale+y_lasso*2)/3  
})  
submission.to_csv('submission3.csv',index=False)
```

Flaws and Next step

Conclusion: After cleaning the data, encoding and model fitting, the project had a 20% root mean squared error.

Flaws: (1) After The PCA deductions, the data retain 35 dimensions, which are still too much. (2) I did not use the features importance selections to choose from original variables. (3) I did not use the stack modelling to combine different models, which means the model might be overfitting.

Next step: (1) Use features importance methods to select models. (2) Explore other ways to deduct the dimensions (3) Use stack models to combine all models

